

UNIX and C Programming (COMP1000)

Lecture 9: Miscellaneous C

Updated: 22nd October, 2019

Department of Computing
Curtin University

Copyright © 2019, Curtin University
CRICOS Provide Code: 00301J

Textbook Reading (Hanly and Koffman)

For more information, see the weekly reading list on Blackboard.

- ▶ **Section 7.7** covers enumerated types.
- ▶ **Section 10.6** covers unions.
- ▶ **Appendix C** covers bitwise operators.

Outline

Timing

Random Numbers

Const

Enumerations

Unions

Bit Manipulation

Timing

- ▶ C has a range of functions and data types for manipulating times and dates.
- ▶ We'll only cover one function here — `time()`.
- ▶ `time()` returns the number of seconds since an arbitrary date.
- ▶ On UNIX, the date is the “Epoch” — 1 January 1970 (midnight UTC) .
- ▶ You can use `time()` to measure an interval of time, or determine the calendar date.

Using time()

- ▶ time() takes one pointer parameter, which you can safely set to NULL.
- ▶ It will return the time in seconds.
- ▶ (If the parameter is not NULL, the time will also be stored at the given address.)
- ▶ The return type is “time_t”, which is just a specially-sized integer.

Example

This code outputs the time taken to execute a function:

```
time_t startTime = time(NULL);  
doStuff();  
printf("%d secs elapsed\n", time(NULL) - startTime);
```

Random Numbers

- ▶ A lot of software relies on randomness or unpredictability (e.g. games, screensavers, password generators, etc.)
- ▶ Randomness is achieved by generating random numbers.
- ▶ In practice, these are usually “pseudorandom” numbers.
- ▶ They’re not truly random, but they’re close enough that nobody will notice.
 - ▶ Based on encryption algorithms.

Random Number Seed — srand()

- ▶ Before you can generate random numbers, you need a “seed”.
- ▶ This is used to obtain a sequence of random numbers.
- ▶ If you always use the same seed, you’ll always get the same sequence of “random” numbers!
- ▶ Set the seed with srand().
- ▶ Takes one int parameter (the seed) and returns void.

Example

Use srand() with the time() function:

```
srand(time(NULL)); /* At start of your program */  
...
```

Generating Random Numbers — rand()

- ▶ Use rand() to obtain random numbers (after srand()).
- ▶ Takes no parameters and returns a random int.
- ▶ Returns a different value each call.

Example

Prints some random numbers:

```
int i;  
srand(time(NULL));  
for(i = 0; i < 4; i++) {  
    printf("%d ", rand());  
}
```

Output: 1034290004 1846603216 1185721333 1246567316

Seed Once (Usually)

- ▶ In most cases, run `srand()` *once only*, when your program starts.
- ▶ Do *not* do this:

```
int i;  
for(i = 0; i < 4; i++) {  
    srand(time(NULL));    /* <-- Very silly! */  
    printf("%d ", rand());  
}
```

You'll get the *same* value repeated 4 times! Why?

- ▶ `srand()` resets the sequence.
- ▶ And we're giving it the same seed each time.
- ▶ (Why? Because if `time()` is called multiple times per second, it must return the same value.)
- ▶ Only run `srand()` more than once if you need to *repeat* a particular sequence of pseudorandom numbers.

Random Numbers Within a Range

- ▶ `rand()` returns values between 0 and `RAND_MAX`.
- ▶ `RAND_MAX` is a platform-specific constant.
- ▶ It is not a very useful limit.
- ▶ You almost always need to constrain the range of numbers.
- ▶ Use division or the modulus (%) operator to do this.

Examples

To get a random int between 0 and `limit - 1`:

```
rand() % limit
```

To get a random double between 0.0 and 1.0:

```
(double)rand() / (double)RAND_MAX
```

Const

- ▶ Until now, you've used `#define` to create constants:

```
#define PI 3.14159
```

- ▶ Done by the preprocessor, through by substitution.
 - ▶ Has no particular datatype.
- ▶ You can also use the `const` keyword:

```
const double PI = 3.14159;
```

- ▶ Done by the C language itself, *not* the preprocessor.
 - ▶ Has a specific datatype.

Const Storage

- ▶ `const` values are (generally) stored in memory, like variables.
- ▶ In memory, a “`const int`” looks just like an “`int`”.
- ▶ The “`const`” keyword is just a pledge that you won’t modify the value.
- ▶ (You *can* modify a “`const`” value, by abusing pointers and typecasting, but don’t!)

Pointers to Constants

- ▶ Since `const` values are stored in memory, you can have pointers to them.
- ▶ This creates some complex situations!

```
const int x = 100;
const int y = 200;
const int * ptr; /* Pointer to an int constant */

ptr = &x;
ptr = &y;
*ptr = 300; /* Invalid (*ptr is constant) */
```

- ▶ Here, you can't modify `x`, `y` or `*ptr` — they're constants.
- ▶ You *can* modify the pointer address `ptr`, making it point somewhere else.

Variables as Constants

- ▶ You can effectively turn a variable into a constant, temporarily:

```
int x = 100;  
const int * ptr = &x; /* Ptr to an int constant */  
  
x = 200;  
*ptr = 300; /* Invalid (*ptr is constant) */
```

- ▶ Here, x is a variable, but *ptr is a constant.
- ▶ They both represent the same location in memory!

Functions and Constants

Functions can accept constant parameters:

```
void func(const int * ptr) {...}
```

- ▶ *ptr is constant *within* the function.
- ▶ We can still pass in an ordinary, non-constant variable:

```
int x = 100; /* Here, x is a normal variable */  
func(&x);  
/* x is guaranteed to still be 100 */
```

- ▶ So, we're only restricting the function. Why bother?
- ▶ When you want to call the function, you instantly know that it's not going to modify the values you pass in.

Constant Pointers

- ▶ Pointers themselves can be constant:

```
int x = 100;
int y = 200;
int *const ptr = &x; /* Constant ptr to an int */

*ptr = 200;
ptr = &y; /* Invalid (ptr is constant) */
```

- ▶ Here, we can modify *ptr freely.
- ▶ ptr itself cannot be changed — it will always point to x.

Constant Pointers to Constants

- You can fix both the pointer and the thing it points to:

```
int x = 100;
int y = 200;

/* Constant pointer to a constant int. */
const int *const ptr = &x;

ptr = &y;      /* Invalid (ptr is constant) */
*ptr = 300;   /* Invalid (*ptr is constant) */
```

Const Declarations

- ▶ Read declarations **right to left**.
- ▶ For example:

```
/* Variable pointers to constant ints. */  
const int * ptr1a;  
int const * ptr1b;  
  
/* Constant pointer to a variable int. */  
int *const ptr2;  
  
/* Constant pointers to constant ints. */  
const int *const ptr3a;  
int const *const ptr3b;
```

(The order of `int` and `const` at the start makes no difference.)

Arrays of Constants

- ▶ You can have an entire array of constants:

```
const int array[] = {5, 10, 15, 20, 25};
```

- ▶ This works like a normal array, but you can't modify its contents.
- ▶ You can also have an array of constant strings:

```
const char *const months[12] = {  
    "January", "February", "March", ...};
```

Global Constants

- ▶ Constants can be local or global, like variables.
- ▶ Global constants are *not* (necessarily) bad practice — they can be useful.

```
const int daysInMonth[12] = {31, 28, 31, ...};  
  
/* These functions all use 'daysInMonth' */  
void printCalendar(int month);  
int isValidDate(int day, int month, int year);  
...
```

Volatile

- ▶ Volatile variables are used in hardware programming.
- ▶ A volatile variable can:
 - ▶ cause side-effects when accessed; and
 - ▶ be modified by something external to the program.
- ▶ Consequently, the compiler cannot perform optimisations
- ▶ Volatile variables use the “volatile” keyword:

```
volatile int a;  
volatile int * b;  
int *volatile c;  
const int *volatile d;  
const volatile int e;  
volatile int *const *volatile f;
```

- ▶ volatile can be used in the same way as const.
- ▶ volatile and const can be used together.

Enumerations

- ▶ An enumeration is a data type whose values are labels or categories.
- ▶ Enumerations are programmer-defined — there can be many different types (like structs).
- ▶ Enumerations only have one value at a time (like `ints`).
- ▶ Enumerations are *not* strings or chars!

Examples

You could define enumeration types to store the following data:

- ▶ A compass direction (north, south, east or west).
- ▶ A menu option (add, delete, search, etc.).

Enumeration Syntax

To define an enumeration:

```
/* Type declaration */  
enum Direction {NORTH, EAST, SOUTH, WEST};  
...  
  
/* Variable declaration */  
enum Direction dir = WEST;
```

OR

```
/* Type declaration */  
typedef enum {NORTH, EAST, SOUTH, WEST} Direction;  
...  
  
/* Variable declaration */  
Direction dir = WEST;
```

Using Enumerations

- ▶ Alternatives that have no natural numeric value.
- ▶ Most operators (+, /, >=, etc) don't make sense for enums.
- ▶ Often used with switch statements:

```
typedef enum {NORTH, EAST, SOUTH, WEST} Direction;  
...  
Direction dir;  
int x, y;  
...  
switch(dir) {  
    case NORTH: y++; break;  
    case EAST:  x++; break;  
    case SOUTH: y--; break;  
    case WEST:  x--; break;  
}
```


Enumeration Examples

```
typedef enum {CLUBS, HEARTS, SPADES, DIAMONDS} Suit;
```

```
typedef enum {  
    OOPD, EP, UCP, ...  
} Unit;
```

```
typedef struct {  
    int numUnits;  
    Unit* units;  
    enum {FULLTIME, PARTTIME} studyMode;  
    enum {  
        GOOD, CONDITIONAL, TERMINATED, GRADUATED  
    } standing;  
} Student;
```

Enumerations as Integer Constants

- ▶ Behind the scenes, enum labels are really just integer constants.
- ▶ By default, the C compiler assigns values 0, 1, 2, etc. to each label in turn.
- ▶ For instance:

```
typedef enum {NORTH, EAST, SOUTH, WEST} Direction;  
  
/* NORTH == 0  
   EAST  == 1  
   SOUTH == 2  
   WEST  == 3 */  
  
...  
printf("%d\n", SOUTH); /* Prints "2" */
```

Choosing Enumeration Values

- ▶ Normally, you don't care about the actual values — just an implementation detail.
- ▶ By default:
 - ▶ The first label is zero.
 - ▶ Each other label is one more than the previous.
- ▶ However, you can choose some or all of them yourself:

```
typedef enum {  
    NORTH, EAST = 3, SOUTH = 10, WEST} Direction;  
  
/* NORTH == 0  
   EAST == 3  
   SOUTH == 10  
   WEST == 11 */
```

(You can only choose integer values.)

Unions

- ▶ A “union” looks like a struct. . .
- ▶ All the fields are stored in the same memory location!
- ▶ Obviously, you can only use one field at a time.
- ▶ Allows you to re-use the same memory for different things.
- ▶ Not used very often.

Union Syntax

You declare and access a union just like a struct:

```
typedef union {  
    int intValue;  
    char strValue[10];  
} Ident;  
  
...  
  
Ident id;  
id.intValue = 42; /* Overwrites id.strId! */
```

- ▶ `&id.intValue == id.strValue`
- ▶ Writing to one will corrupt the other!

Union Examples

```
typedef union {  
    int *intp;  
    double *doublep;  
    char *charp;  
} MultiPointer;
```

```
typedef struct {  
    enum {CIRCLE, RECTANGLE} type;  
    union {  
        double radius;  
        struct {float width; float height;} rect;  
    } dimensions;  
} Shape;
```

Bit Manipulation

- ▶ All data is (of course) an array of bits.
- ▶ However, bits cannot be read or written to directly.
- ▶ Here we discuss some tricks to manipulate them.

Visualisation

This is binary form of the 8-bit char value 'Z':

7	6	5	4	3	2	1	bit 0
0	1	0	1	1	0	1	0

(ints, doubles, etc. are much longer, of course.)

Shift Left and Shift Right

- ▶ << is the shift-left operator.
- ▶ >> is the shift-right operator.
- ▶ These shift a bit pattern left or right, by a given amount.
- ▶ For ints, this effectively multiplies or divides by powers of two.


Example

```
int var = 100;


printf("%d\n", var << 1); /* 200 */
printf("%d\n", var << 2); /* 400 */
printf("%d\n", var << 3); /* 800 */

printf("%d\n", var >> 1); /* 50 */
printf("%d\n", var >> 2); /* 25 */
```



Shift Left — Graphical Example




Shift Left — Graphical Example




Shift Left — Graphical Example




Shift Left — Graphical Example




Shift Left — Graphical Example




Shift Left — Graphical Example



Shift Left — Graphical Example



Shift Left — Graphical Example



Bitwise Operators

- ▶ Remember the “logical” operators: `&&`, `||` and `!`.
- ▶ The “bitwise” operators are: `&`, `|`, `^` and `~`.
- ▶ Logical and bitwise operators do the same thing, except:
 - ▶ Logical operators work on whole variables.
 - ▶ Bitwise operators work on each bit separately.
- ▶ With n -bit data, bitwise operators will perform n separate AND/OR/XOR/NOT operations.
- ▶ The n resulting bits are returned as one value.

Bitwise NOT (~)

- ▶ The simplest bitwise operator.
- ▶ A unary operator (like “!”) — takes only one operand.
- ▶ Inverts each bit in the operand — ones become zeros, zeros become ones.

7	6	5	4	3	2	1	bit 0
0	1	0	1	1	0	1	0

`char ch = 'Z'; (90)`

1	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

`~ch (165)`

Bitwise AND (&)

- ▶ Same symbol as the address-of operator, but different operation.
- ▶ “a & b” calculates the bitwise AND of a and b.
- ▶ Each resulting bit is 1 if *both* operand bits are 1.

7	6	5	4	3	2	1	bit 0
0	1	0	1	1	0	1	0

 char z = 'Z'; (90)

bitwise-AND

0	1	0	0	0	1	1	0
---	---	---	---	---	---	---	---

 char f = 'F'; (70)

=

0	1	0	0	0	0	1	0
---	---	---	---	---	---	---	---

 z & f (66, 'B')

Bitwise OR (|)

- ▶ “a | b” calculates the bitwise OR of a and b.
- ▶ Each resulting bit is 1 if *either* operand bit is 1.

7	6	5	4	3	2	1	bit 0	
0	1	0	1	1	0	1	0	char z = 'Z'; (90)

bitwise-OR

0	1	0	0	0	1	1	0	char f = 'F'; (70)
---	---	---	---	---	---	---	---	--------------------

=

0	1	0	1	1	1	1	0	z f (94)
---	---	---	---	---	---	---	---	------------

Bitwise XOR (^)

- ▶ XOR means “exclusive OR”.
- ▶ “ $a \wedge b$ ” calculates the bitwise XOR of a and b .
- ▶ Each resulting bit is 1 if the two operand bits are *different* (i.e. only one of the operands is 1).
- ▶ Used in encryption and hash functions.

7	6	5	4	3	2	1	bit 0	
0	1	0	1	1	0	1	0	<code>char z = 'Z'; (90)</code>
bitwise-XOR								
0	1	0	0	0	1	1	0	<code>char f = 'F'; (70)</code>
=								
0	0	0	1	1	1	0	0	<code>z ^ f (28)</code>

Bit Fields

- ▶ You can place “bit fields” inside structs (and unions).
- ▶ These are integers that occupy a specified number of bits.
- ▶ They can be signed (allowing for negative numbers) or unsigned.
- ▶ The number of bits determines how large their values can be:
 - ▶ 1 bit allows only zero and one.
 - ▶ 2 bits allows values 0 to 3 (unsigned) or -2 to 1 (signed).
 - ▶ 3 bits allows values 0 to 7 (unsigned) or -4 to 3 (signed).
 - ▶ 4 bits allows values 0 to 15 (unsigned) or -8 to 7 (signed).
 - ▶ etc.
- ▶ The syntax is:

```
unsigned int <name> : <bits>;
```

OR

```
signed int <name> : <bits>;
```

Bit Fields — Example

```
typedef struct {  
    unsigned int flag : 1;           /* 0 or 1 */  
    unsigned int quadrant : 2;       /* 0 to 3 */  
    signed int rating : 4;           /* -8 to 7 */  
    int normalInt;                   /* much larger */  
} SmallThings;
```

Here, the SmallThings struct has 3 bit fields and one normal int field.

Bit Fields — Limitations

- ▶ Unfortunately, pointers can only point to bytes, not bits.
- ▶ Therefore, you can't have a pointer to a bit field.
- ▶ The address-of operator won't work on bit fields.
- ▶ You can't pass a bit field by reference.
- ▶ You can't use it with `scanf()` (but `printf()` will work).

Coming Up

- ▶ That's the last of the C lectures!
- ▶ In the final lecture next week, you'll get a taste of C++.